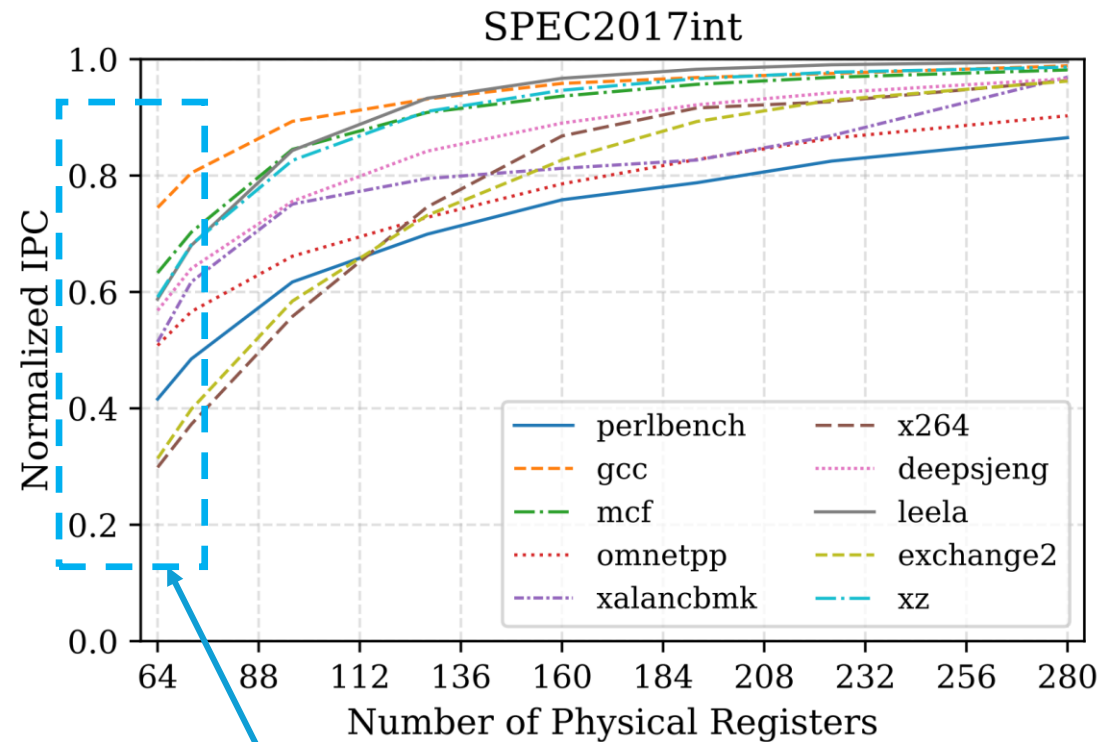
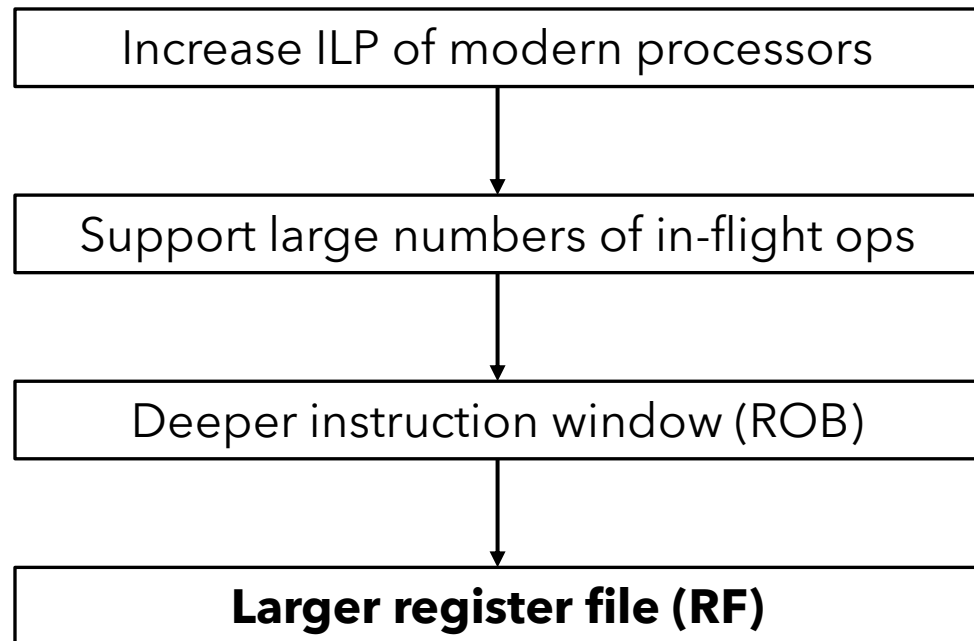


ATR: Out-of-Order Register Release Exploiting Atomic Regions

Yinyuan Zhao, Surim Oh, Mingsheng Xu, and Heiner Litz
University of California, Santa Cruz

Scaling RF Size is Important



With 64 physical registers, the average IPC reaches only 37.7% of the ideal case

A larger RF is crucial for sustaining high IPC

Scaling RF Size is Challenging


10-wide u-Arch

- 20 read RF ports
- 10 write RF ports

+

512-entry RF

- 512:1 multiplexers
- 64-bit flip-flops



2M transistors

$$\begin{array}{r} 1024 \times 64 \times 20 \\ + \quad 512 \times 64 \times 20 \\ \hline 1.3M + 650K \end{array}$$

Introduce challenges in area, power, and clock frequency

Conventional Renaming

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed

$R \leftarrow r_4, r_5$

I3 Redefined

I3 Committed

Conventional Renaming

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* is not used in the future

$R \leftarrow r_4, r_5$

I3 Redefined

I3 Committed

Conventional Renaming

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* is not used in the future

$R \leftarrow r_4, r_5$

I3 Redefined

I3 Committed



The *ptag* of register R is released

Conventional Renaming holds physical registers much longer than needed!

Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* of register R is **early released**

----- **Branch**

$R \leftarrow r_4, r_5$

I3 Redefined

Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* of register R is **early released**

----- ***Branch***

$R \leftarrow r_4, r_5$

I3 Redefined



The redefined instruction is mispredicted

Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* of register R is **early released**

----- **Branch**

$R \leftarrow r_4, r_5$

I3 Redefined



The redefined instruction is mispredicted

$r_2 \leftarrow R, r_3$

I4 Consumed



The released *ptag* is consumed!!!

Speculative early release is unsafe or not feasible

Non-Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* is not used in the future

----- ***Branch***

$R \leftarrow r_4, r_5$

I3 Redefined

Non-Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* is not used in the future

----- **Branch**

$R \leftarrow r_4, r_5$

I3 Redefined

All the older branches or exception-causing instructions are executed

I3 Precommitted

Non-Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* is not used in the future

Branch

$R \leftarrow r_4, r_5$

I3 Redefined

All the older branches or exception-causing instructions are executed

I3 Precommitted



The *ptag* is safely early released

Non-Speculative Early Release

$R \leftarrow r_1, r_2$

I1 Renamed



Register R is renamed by allocating a *ptag*

$r_2 \leftarrow R, r_3$

I2 Consumed



The *ptag* is not used in the future

Unused

Branch

$R \leftarrow r_4, r_5$

I3 Redefined

All the older branches or exception-causing instructions are executed

I3 Precommitted



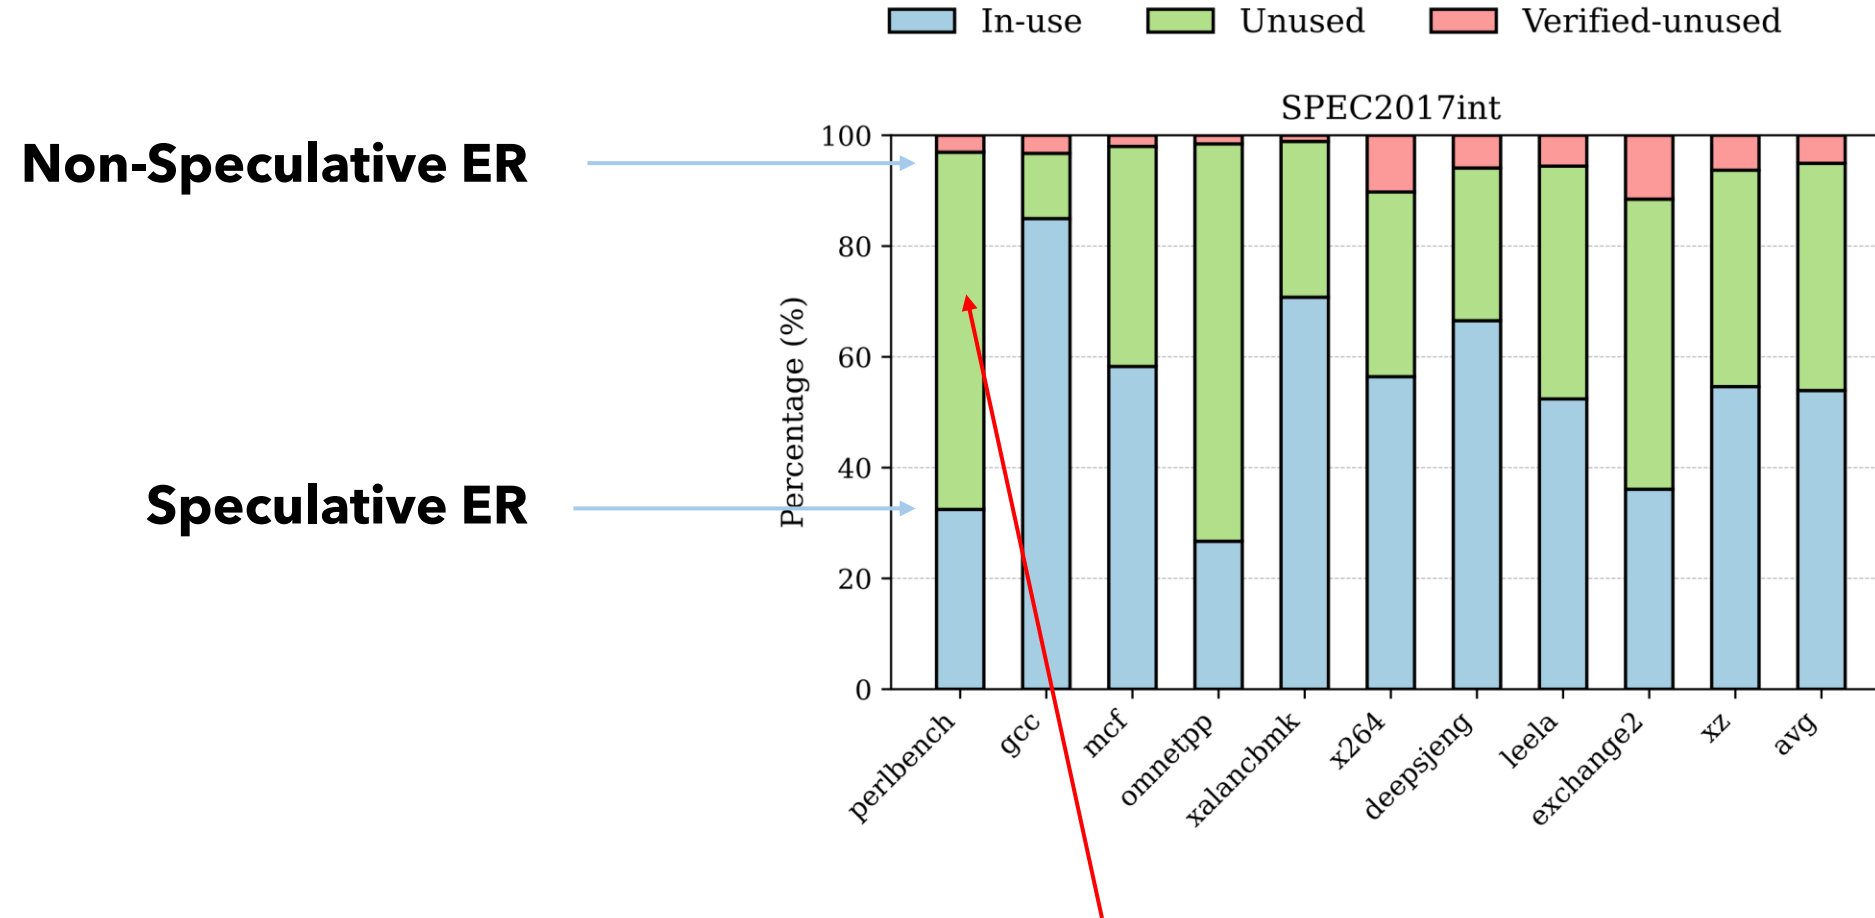
The *ptag* is safely early released

**Verified-
Unused**

I3 Committed

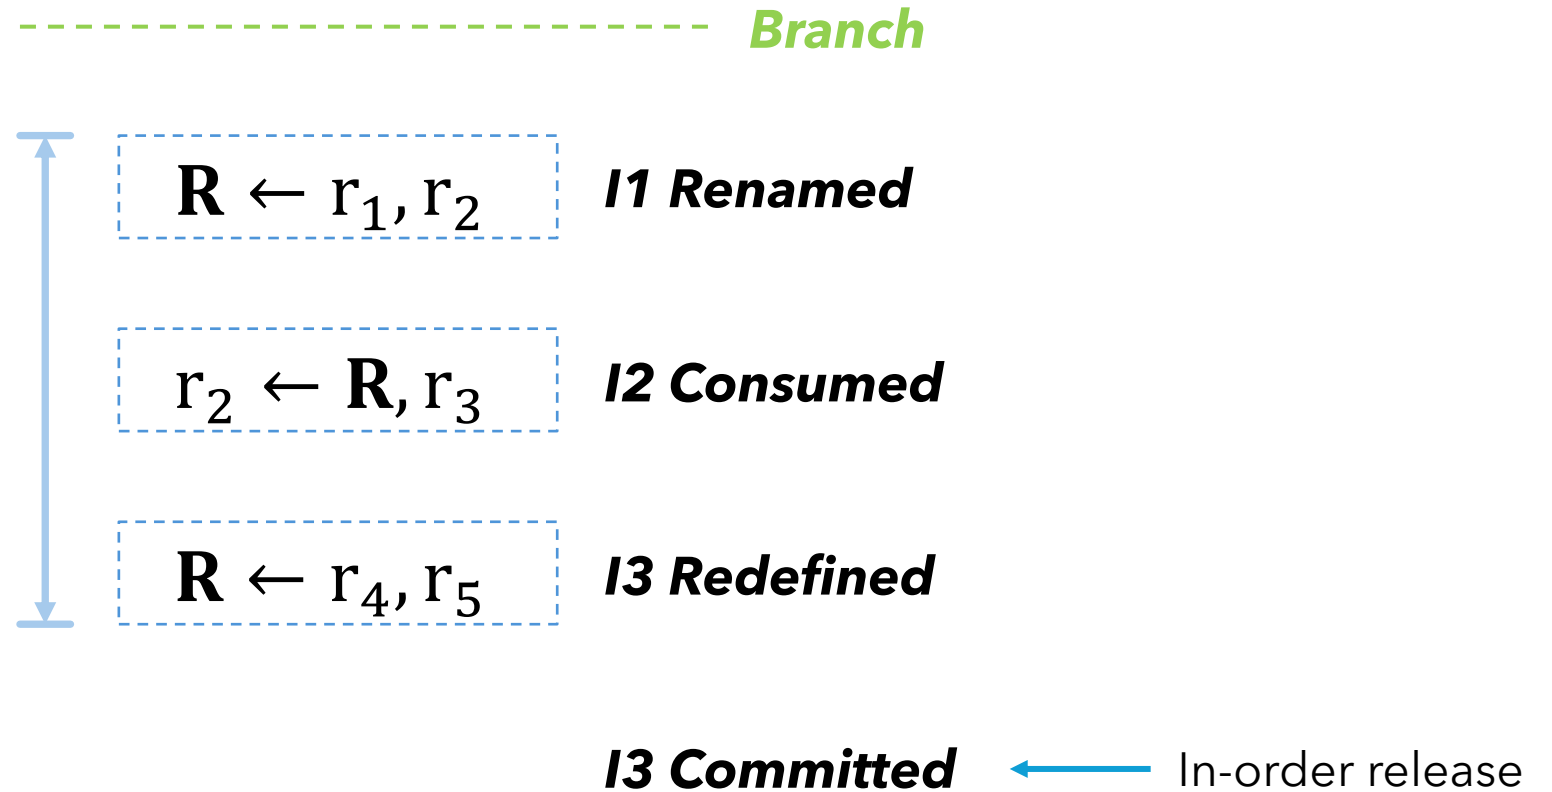
Conv.

Gap between Non-Spec and Spec ER



Opportunity of early releasing unused register

Our Proposal

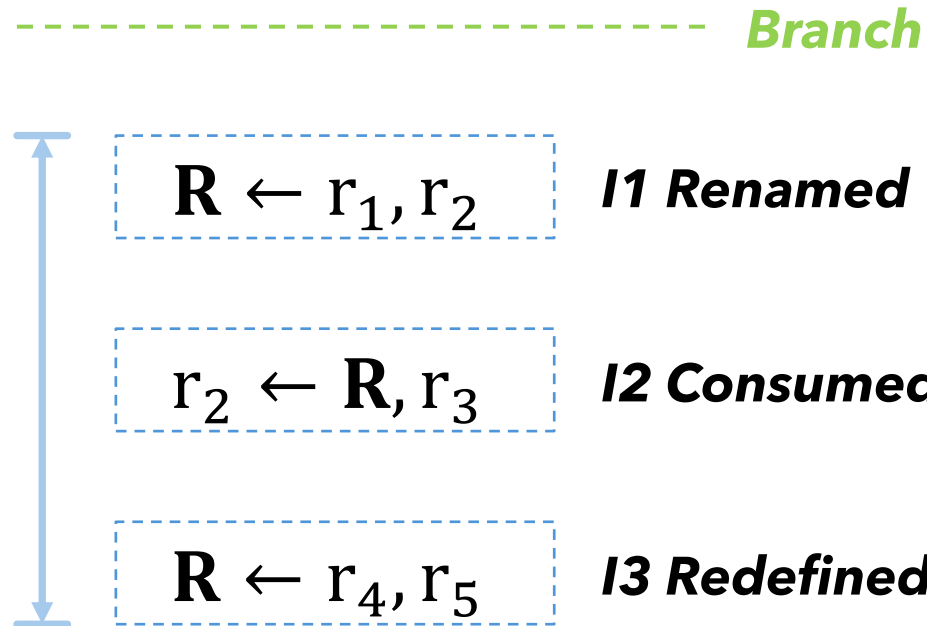


Our Proposal

Atomic Commit Region

Branches/Exceptions-free

- Either entirely off-path or on-path



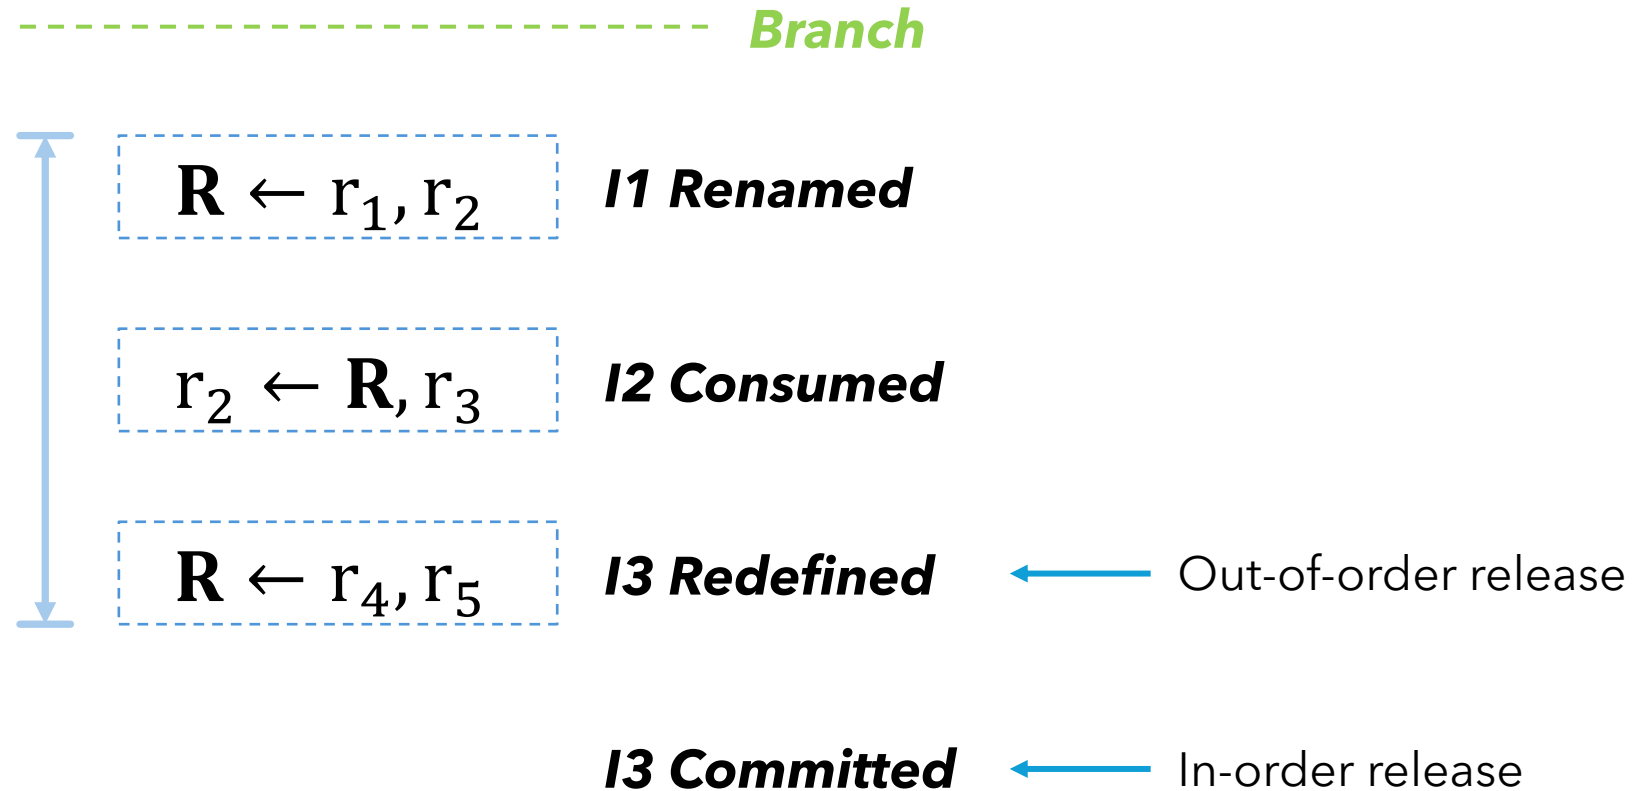
I3 Committed ← In-order release

Our Proposal

Atomic Commit Region

Branches/Exceptions-free

- Either entirely off-path or on-path



Does not violate precise exceptions nor suffers from mispredictions

Analysis

			<u>Re</u>	<u>Ex</u>	<u>Cm</u>	<u>Pr</u>
I1	MOVE	RAX \leftarrow RAX	510	675	839	675
I2	TEST + JNZ	ZPS \leftarrow RAX	510	841	841	841
				
I3	LEA	RAX \leftarrow RDI	709	716	716	841
				
I4	LEA	RBX \leftarrow RAX	729	737	737	842
I5	SHR	RBX \leftarrow RBX, ZPS	729	738	738	842

← I1 is a high-latency load

*An example segment of instructions from
SPEC2017int omnetpp*

Analysis

			<u>Re</u>	<u>Ex</u>	<u>Cm</u>	<u>Pr</u>
I1	MOVE	RAX \leftarrow RAX	510	675	839	675
I2	TEST + JNZ	ZPS \leftarrow RAX	510	841	841	841
...				
I3	LEA	RAX \leftarrow RDI	709	716	716	841
...				
I4	LEA	RBX \leftarrow RAX	729	737	737	842
I5	SHR	RBX \leftarrow RBX, ZPS	729	738	738	842



I1 is a high-latency load



I2 needs to wait until wake-up

*An example segment of instructions from
SPEC2017int omnetpp*


Analysis

			<u>Re</u>	<u>Ex</u>	<u>Cm</u>	<u>Pr</u>	
I1	MOVE	RAX ← RAX	510	675	839	675	← I1 is a high-latency load
I2	TEST + JNZ	ZPS ← RAX	510	841	841	841	← I2 needs to wait until wake-up
...					
I3	LEA	RAX ← RDI	709	716	716	841	
...					
I4	LEA	RBX ← RAX	729	737	737	842	← I4 precommit is delayed
I5	SHR	RBX ← RBX, ZPS	729	738	738	842	

*An example segment of instructions from
SPEC2017int omnetpp*


Lots of cycles are wasted in waiting for precommitting

Atomic Ratio



From the atomic register to the redefined one:


- no branch/jump
- no exception-causing instructions (ld/st/div)



Over 17% of all allocated registers are located within atomic commit regions

ATR Algorithm

- **1. Atomic**
No **br/jmp/ld/st/div** between the renaming and redefining instruction



- **2. Consumed**
All consumers have consumed R
- **3. Redefined**
R has been redefined by a later instruction

No need to wait for commit or pre-commit of any instruction!

Design

Consumer Renaming

- Consumer Count ++

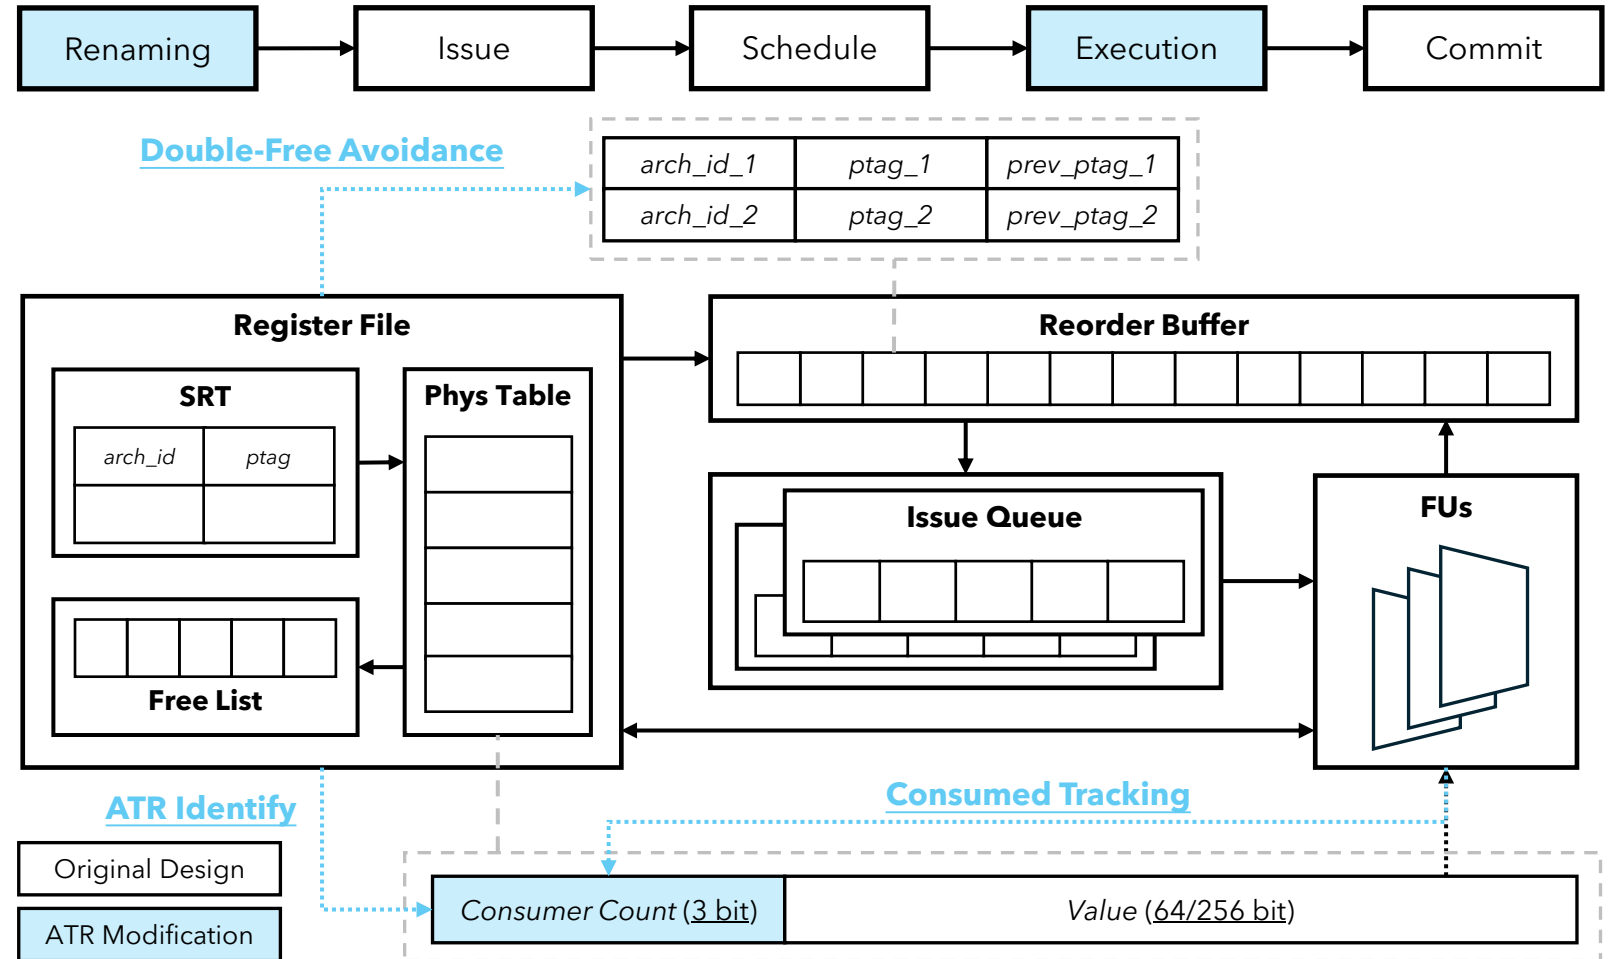
Consumer Execution

- Consumer Count --

*3-bit consumer count
for each ptag*

BR/JMP/LD/ST/DIV Renaming

- Consumer Count = MAX



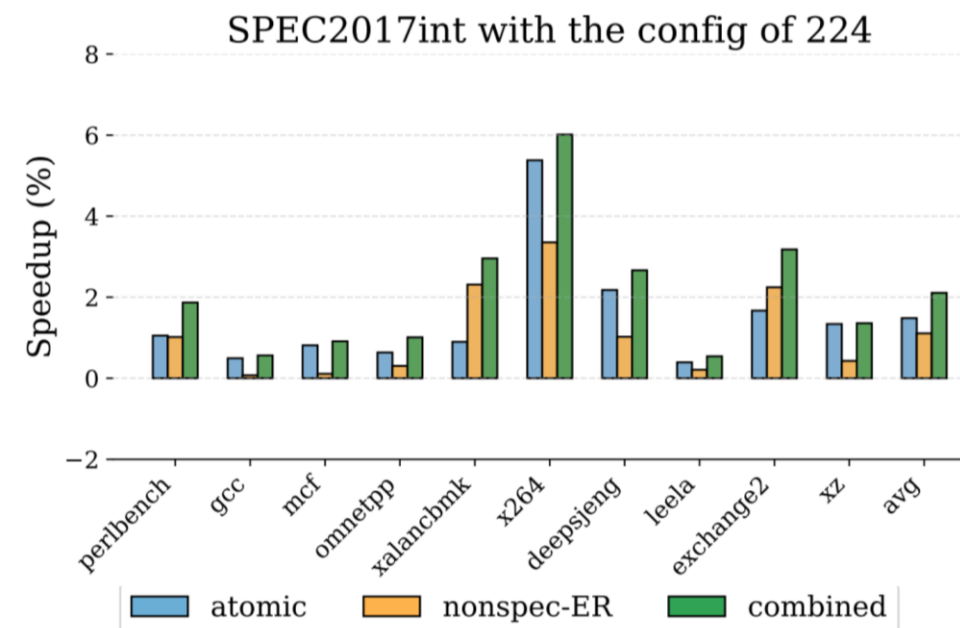
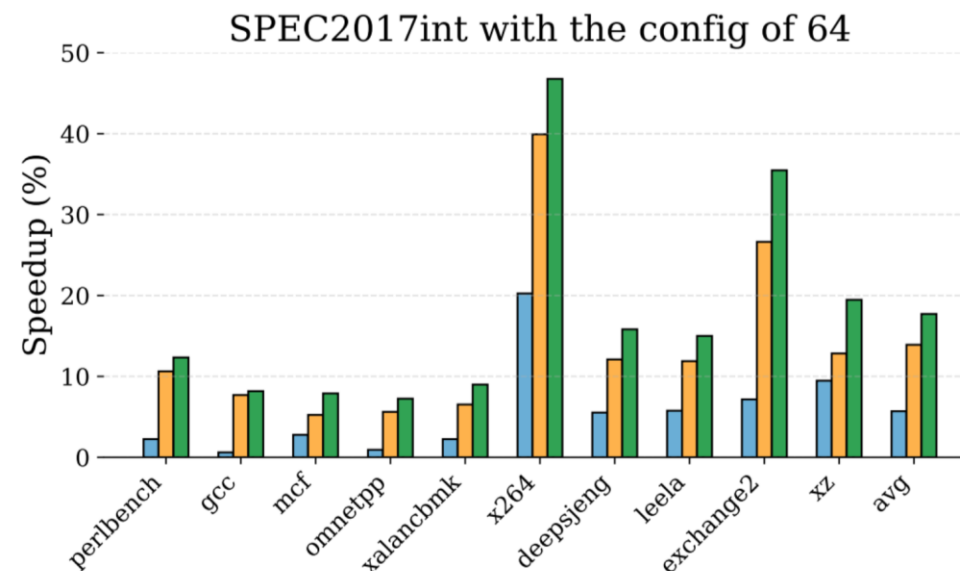
Evaluation

Configuration

- Arch: Golden_Cove
- Dataset: SPEC2017int

Results

- @64
 - atomic: + 5.70% over baseline
 - combine: + 3.23% over nonspec
- @224
 - atomic: + 1.48% over baseline
 - + 0.37% over nonspec



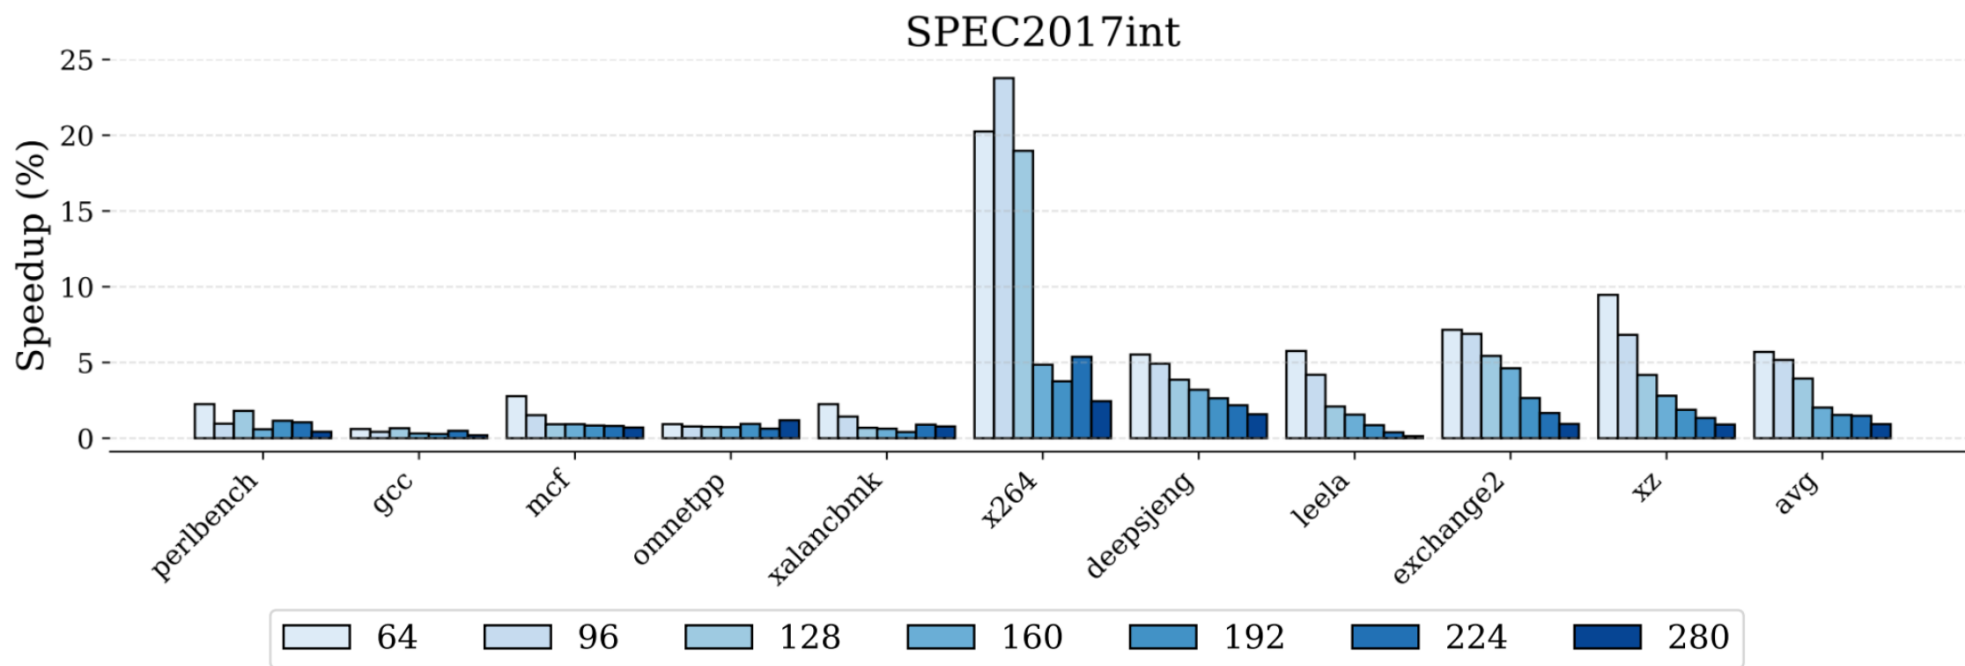
Evaluation

Configuration

- Arch: Golden_Cove
- Dataset: SPEC2017int

Results

- The speedup decreases as the register size increases



Thank you!